

Automatic Differentiation

Lecture 18b

What is “Automatic Differentiation”

- We want to compute both $f(x)$, $f'(x)$, given a program for $f(x) := \exp(x^2) * \tan(x) + \log(x)$.
- We could do this via...

$$\frac{d \left(e^{x^2} \tan x + \log x \right)}{dx} = 2x e^{x^2} \tan x + e^{x^2} \sec^2 x + \frac{1}{x}$$

But this is too much work.

- All we really want is a way to compute, for a specific numeric value of x , $f(x)$, and $f'(x)$. Say the value is $x=2.0$. We could evaluate those two expressions, or we could evaluate a Taylor series for f around the point 2.0, whose coefficients would give us $f(2)$ and $f'(2)/2!$
- Or we could do something cleverer.

ADOL, ADOL-C, Fortran (77, 90) or C “program differentiation”

- Rewrite (via a compiler?) a more-or-less arbitrary Fortran program that computes $f(c)$ into a program that computes $\langle f(c), f'(c) \rangle$ for a real value c .
- How to do this?
- Two ways, “forward” and “reverse”

Forward automatic differentiation

- In the program to compute $f(c)$ make these substitutions:
 - Where we see the expression x , replace it with $\langle c, 1 \rangle$
 - the rule $d(u) = 1$ if $u = x$.
 - Where we see a constant, e.g. 43, use $\langle 43, 0 \rangle$
 - the rule $d(u) = 0$ if u does not depend on x
 - Where we see an operation $\langle a, b \rangle + \langle r, s \rangle$ use $\langle a+r, b+s \rangle$
 - the rule $d(u+v) = du+dv$
 - Where we see an operation $\langle a, b \rangle * \langle r, s \rangle$ use $\langle a*r, b*r+a*s \rangle$
 - the rule $d(u*v) = u*dv+v*du$
 - Where we see $\cos(\langle r, s \rangle)$ use $\langle \cos(r), -\sin(r)*s \rangle$
 - the rule $d(\cos(u)) = -\sin(u)du$
 - Etc.

Two ways to do forward AD

- Rewrite the program so that every variable is now TWO variables, var and var_x , the extra var_x is the derivative.
- Leave the “program” alone and overlay the meaning of all the operations like $+$, $*$, \cos , with generalizations on pairs $\langle \alpha, \beta \rangle$

Does this work “hands off”? (Either method..)

- Mostly, but not entirely.
- Some constructions are not differentiable. Floor, Ceiling, decision points.
- Some constructions don't exist in normal fortran (etc.) e.g. “get the derivative” needed for use by
 - Newton iteration: $z_{i+1} := z_i - f(z_i)/f'(z_i)$
- Is it still useful? Apparently enough to build up a minor industry in enhanced compilers. see www.autodiff.org for comprehensive references
- If the programs start in Lisp, the transformations are clean and short. See www.cs.berkeley.edu/~fateman/papers/overload-AD.pdf

Where is this useful?

- A number of numerical methods benefit from having derivatives conveniently available: minimization, root-finding, ODE solving.
- Generalized to $F(x,y,z)$, and partial derivatives of order 2, 3, (or more?)

What about reverse differentiation?

- The reverse mode computes derivatives of dependent variables with respect to intermediate variables and propagates from one statement to the previous statement according to the chain rule.
- Thus, the reverse mode requires a “reversal” of the program execution, or alternatively, a stacking of intermediate values for later access.

Reverse..

- Given independent vars $\mathbf{x} = \{x_1, x_2, \dots\}$
- And dependent vars $\mathbf{y} = \{y_1(\mathbf{x}), y_2(\mathbf{x}) \dots\}$, and also temporary vars...
- (read grey stuff upward...)
- $A = f(x_1, x_2)$
- $B = g(A, x_3)$... we know d/dx_1 of B is $dG/d1 * dA/dx_1$
- $R = A + B$... we know d/dx_1 of R is $dA/dx_1 + dB/dx_1$, what are they?

Which is better?

- Hybrid methods have been advocated to take advantage of both approaches. Reverse has a major time advantage for many vars; it has the disadvantage of non-linear memory use—depends on program flow, loops iterations must be stacked.
- Numerous conferences on this subject.
- Lots of tools, substantial cleverness.